

Matthew Morgan
Student ID: 010471280
Data Mining II - 212
Task 2: Dimensionality Reduction Methods
Western Governor's University
Program Mentor: Mandy Rasmuson

Part I: Research Question

A. Describe the purpose of this data mining report by doing the following:

A1. Propose one question relevant to a real-world organizational situation that you will answer by using principal component analysis (PCA).

The purpose of this task is to try and use dimensionality reduction methods to help answer a research question. My research question for this project will be if we can help determine if patient re-admissions can be predicted through the use of a decision tree classifier after using principal component analysis (PCA) to reduce the dimensionality of the provided dataset.

A2. Define one goal of the data analysis. Ensure that your goal is reasonable within the scope of the scenario and is represented in the available data.

The ultimate goal of this performance assessment is to predict patient re-admissions. Based on the data dictionary this is the primary goal of the WGU Health Data as high re-admission rates result in fines. If we can find out how to reduce patient re-admissions not only will we be able to save money by avoiding fines we can then use that money towards prevention and provide better health outcomes for the patients.

Part II: Method Justification

B. Explain the reasons for using PCA by doing the following:

B1. Explain how PCA analyzes the selected data set. Include expected outcomes.

PCA is a way we can reduce dimensions within a dataset. This can be done through feature elimination or extraction. Elimination is just that, it removes features that are redundant or have no correlation to the dependent variable. Extraction involves combining multiple features into one datapoint. This new datapoint is then evaluated to determine how important it is by how it effects the variance within the dataset. If it meets the threshold that is defined it is then considered a "principal component" in the appropriate model.

The expected outcome is to hopefully find a number of principal components that allow us to predict patient re-admissions. This will be evaluated by using the `explained_variance_ratio_` attribute. The ones that are most important to the model will be saved, and the other eliminated so we can reduce the dimensionality of the dataset.

B2. Summarize one assumption of PCA.

One assumption of PCA provided by Datacamp is that, "the principal components having the highest variance are more important than those which don't." (Keita, 2023) Because we'll be taking the variance scores of the different principal components we will be guided towards what to keep and what we can eliminate based on that assumption.

Part III: Data Preparation

C. Perform data preparation for the chosen dataset by doing the following:

C1. Identify the continuous dataset variables that you will need in order to answer the PCA question proposed in part A1.

Variable # Continuous Variable Dictionary

- 1 Latitude
- 2 Longitude
- 3 Population
- 4 Number of Children in Household
- 5 Age
- 6 Income of Patient
- 7 Patient's Vit D Level
- 8 Number of Physician Visits
- 9 Number of full meals eaten
- 9 Number of Vit D Supplements
- 10 Length of initial stay
- 11 Average daily amount charged for hospitalization
- 12 Average daily amount charged for additional charges
- 13 Initial_Admin_Emergency_admission

C2. Standardize the continuous dataset variables identified in part C1. Include a copy of the cleaned dataset.

In [157]: *#DataCleaning*

```
#Import Packages needed for data cleaning and PCA
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import accuracy_score

#Loading the CSV of the default dataset, index_col to prevent duplicated column
df = pd.read_csv(r'C:\Users\mmorg\WGU\D212\medical_clean.csv', index_col=0)
#Get overview of dataset such as # of columns, names, and size
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 10000 entries, 1 to 10000
```

```
Data columns (total 49 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|--------------------|----------------|---------|
| 0 | Customer_id | 10000 non-null | object |
| 1 | Interaction | 10000 non-null | object |
| 2 | UID | 10000 non-null | object |
| 3 | City | 10000 non-null | object |
| 4 | State | 10000 non-null | object |
| 5 | County | 10000 non-null | object |
| 6 | Zip | 10000 non-null | int64 |
| 7 | Lat | 10000 non-null | float64 |
| 8 | Lng | 10000 non-null | float64 |
| 9 | Population | 10000 non-null | int64 |
| 10 | Area | 10000 non-null | object |
| 11 | TimeZone | 10000 non-null | object |
| 12 | Job | 10000 non-null | object |
| 13 | Children | 10000 non-null | int64 |
| 14 | Age | 10000 non-null | int64 |
| 15 | Income | 10000 non-null | float64 |
| 16 | Marital | 10000 non-null | object |
| 17 | Gender | 10000 non-null | object |
| 18 | ReAdmis | 10000 non-null | object |
| 19 | VitD_levels | 10000 non-null | float64 |
| 20 | Doc_visits | 10000 non-null | int64 |
| 21 | Full_meals_eaten | 10000 non-null | int64 |
| 22 | vitD_supp | 10000 non-null | int64 |
| 23 | Soft_drink | 10000 non-null | object |
| 24 | Initial_admin | 10000 non-null | object |
| 25 | HighBlood | 10000 non-null | object |
| 26 | Stroke | 10000 non-null | object |
| 27 | Complication_risk | 10000 non-null | object |
| 28 | Overweight | 10000 non-null | object |
| 29 | Arthritis | 10000 non-null | object |
| 30 | Diabetes | 10000 non-null | object |
| 31 | Hyperlipidemia | 10000 non-null | object |
| 32 | BackPain | 10000 non-null | object |
| 33 | Anxiety | 10000 non-null | object |
| 34 | Allergic_rhinitis | 10000 non-null | object |
| 35 | Reflux_esophagitis | 10000 non-null | object |
| 36 | Asthma | 10000 non-null | object |
| 37 | Services | 10000 non-null | object |

```
38 Initial_days      10000 non-null float64
39 TotalCharge       10000 non-null float64
40 Additional_charges 10000 non-null float64
41 Item1             10000 non-null int64
42 Item2             10000 non-null int64
43 Item3             10000 non-null int64
44 Item4             10000 non-null int64
45 Item5             10000 non-null int64
46 Item6             10000 non-null int64
47 Item7             10000 non-null int64
48 Item8             10000 non-null int64
dtypes: float64(7), int64(15), object(27)
memory usage: 3.8+ MB
```

```
In [158]: # More exploration to spot any potential problems with the dataset and what might need cleaned up
pd.set_option("display.max_columns", None)
df.head(5)
```

Out[158]:

| | Customer_id | Interaction | UID | City | State | County | Zip | Lat | Lng | Po |
|-----------|-------------|--------------------------------------|----------------------------------|--------------|-------|--------------|-------|----------|-----------|----|
| CaseOrder | | | | | | | | | | |
| 1 | C412403 | 8cd49b13-f45a-4b47-a2bd-173ffa932c2f | 3a83ddb66e2ae73798bdf1d705dc0932 | Eva | AL | Morgan | 35621 | 34.34960 | -86.72508 | |
| 2 | Z919181 | d2450b70-0337-4406-bdbb-bc1037f1734c | 176354c5eef714957d486009feabf195 | Marianna | FL | Jackson | 32446 | 30.84513 | -85.22907 | |
| 3 | F995323 | a2057123-abf5-4a2c-abad-8ffe33512562 | e19a0fa00aeda885b8a436757e889bc9 | Sioux Falls | SD | Minnehaha | 57110 | 43.54321 | -96.63772 | |
| 4 | A879973 | 1dec528d-eb34-4079-adce-0d7a40e82205 | cd17d7b6d152cb6f23957346d11c3f07 | New Richland | MN | Waseca | 56072 | 43.89744 | -93.51479 | |
| 5 | C544523 | 5885f56b-d6da-43a3-8760-83583af94266 | d2f0425877b10ed6bb381f3e2579424a | West Point | VA | King William | 23181 | 37.59894 | -76.88958 | |

In [159]: *#Because we are only using continuous variables, not much cleaning needs to be done and we can remove columns that are categorical or don't apply to the PCA*

```
# Assign all continuous variables for PCA to X variable
X = df[["Lat", "Lng", "Population", "Children", "Age", "Income", "VitD_levels", "Doc_visits",
        "Full_meals_eaten", "vitD_supp", "Initial_days", "TotalCharge", "Additional_charges"]].copy()
# Define list of column headers
X_columns = list(X.columns)
# Assign patient re-admissions column to y variable
y = df["ReAdmis"]
```

In [160]: *# Standardize X by instantiating the StandardScaler(), then fitting and transforming to X*

```
X_stand = StandardScaler().fit_transform(df[["Lat", "Lng", "Population", "Children", "Age", "Income", "VitD_levels",
        "Full_meals_eaten", "vitD_supp", "Initial_days", "TotalCharge", "Additional_charges"]].copy())
# New dataframe with standardized values for verification
X_stand_df = pd.DataFrame(X_stand, columns=X_columns)
```

In [161]: *#Verifying the means of each column, should be 0 to confirm standardization*

```
X_stand_means = X_stand_df.agg(['mean']).round(2)

print(X_std_means)
```

| | Lat | Lng | Population | Children | Age | Income | VitD_levels | Doc_visits | \ |
|------|------|-----|------------|----------|-----|--------|-------------|------------|---|
| mean | -0.0 | 0.0 | -0.0 | -0.0 | 0.0 | 0.0 | -0.0 | 0.0 | |

| | Full_meals_eaten | vitD_supp | Initial_days | Additional_charges |
|------|------------------|-----------|--------------|--------------------|
| mean | 0.0 | -0.0 | -0.0 | -0.0 |

```
In [162]: #Verifying the standard deviation of each column, should be 1 to confirm standardization
X_stand_stdv = X_stand_df.agg(['std']).round(2)

print(X_stand_stdv)
```

```
      Lat  Lng  Population  Children  Age  Income  VitD_levels  Doc_visits  \
std  1.0  1.0          1.0          1.0  1.0    1.0          1.0          1.0

      Full_meals_eaten  vitD_supp  Initial_days  TotalCharge  \
std                1.0          1.0          1.0          1.0

      Additional_charges
std                1.0
```

```
In [163]: #Checking correlations to determine redundant components
X_stand_df.corr()
```

Out[163]:

| | Lat | Lng | Population | Children | Age | Income | VitD_levels | Doc_visits | Full_meals_eaten | vitD_sup |
|--------------------|-----------|-----------|------------|-----------|-----------|-----------|-------------|------------|------------------|-----------|
| Lat | 1.000000 | -0.112348 | -0.207572 | 0.006373 | -0.007270 | -0.019369 | 0.001493 | 0.008380 | 0.003401 | 0.001285 |
| Lng | -0.112348 | 1.000000 | -0.031979 | -0.014114 | 0.007493 | -0.006665 | -0.006389 | 0.000754 | -0.014231 | -0.001961 |
| Population | -0.207572 | -0.031979 | 1.000000 | 0.002462 | -0.018987 | 0.005426 | 0.002651 | 0.012646 | -0.025608 | 0.009781 |
| Children | 0.006373 | -0.014114 | 0.002462 | 1.000000 | 0.009836 | 0.007176 | 0.009487 | -0.002292 | 0.003835 | -0.004319 |
| Age | -0.007270 | 0.007493 | -0.018987 | 0.009836 | 1.000000 | -0.012228 | 0.010315 | 0.006898 | 0.008555 | 0.010014 |
| Income | -0.019369 | -0.006665 | 0.005426 | 0.007176 | -0.012228 | 1.000000 | -0.013115 | 0.013464 | -0.011365 | 0.001285 |
| VitD_levels | 0.001493 | -0.006389 | 0.002651 | 0.009487 | 0.010315 | -0.013115 | 1.000000 | 0.010210 | 0.023223 | -0.007200 |
| Doc_visits | 0.008380 | 0.000754 | 0.012646 | -0.002292 | 0.006898 | 0.013464 | 0.010210 | 1.000000 | -0.002767 | 0.005681 |
| Full_meals_eaten | 0.003401 | -0.014231 | -0.025608 | 0.003835 | 0.008555 | -0.011365 | 0.023223 | -0.002767 | 1.000000 | -0.019980 |
| vitD_supp | 0.001285 | -0.001961 | 0.009781 | -0.004319 | 0.010014 | 0.001253 | -0.007203 | 0.005681 | -0.019980 | 1.000000 |
| Initial_days | -0.008820 | -0.009292 | 0.017469 | 0.022467 | 0.016264 | -0.012465 | -0.003642 | -0.006754 | -0.017267 | 0.015972 |
| TotalCharge | -0.010759 | -0.008830 | 0.019188 | 0.024100 | 0.016876 | -0.014345 | -0.001403 | -0.005043 | -0.014306 | 0.016922 |
| Additional_charges | -0.002283 | 0.000079 | -0.004820 | 0.013548 | 0.716854 | -0.009825 | 0.008290 | 0.008072 | 0.018763 | 0.010322 |


```
In [164]: #As shown Initial_days and TotalCharge are highly correlated and redundant.
#As I have done in previous classes, I dropped TotalCharge as they both basically represent the same data.
#Longer stays lead to more initial charges.

#Re-creating X variable and X_columns with this in mind
X = df[["Lat", "Lng", "Population", "Children", "Age", "Income", "VitD_levels", "Doc_visits",
        "Full_meals_eaten", "vitD_supp", "Initial_days", "Additional_charges"]].copy()
# Define list of column headers
X_columns = list(X.columns)
# Standardize X by instantiating the StandardScaler(), then fitting and transforming to X
X_stand = StandardScaler().fit_transform(df[["Lat", "Lng", "Population", "Children", "Age", "Income", "VitD_levels",
        "Full_meals_eaten", "vitD_supp", "Initial_days", "Additional_charges"]].copy())
# New dataframe with standardized values for verification
X_stand_df = pd.DataFrame(X_stand, columns=X_columns)
```

```
In [165]: #Verifying the means of each column, should be 0 to confirm standardization
X_stand_means = X_stand_df.agg(['mean']).round(2)

print(X_stand_means)
```

| | Lat | Lng | Population | Children | Age | Income | VitD_levels | Doc_visits | \ |
|------|------|-----|------------|----------|-----|--------|-------------|------------|---|
| mean | -0.0 | 0.0 | -0.0 | -0.0 | 0.0 | 0.0 | -0.0 | 0.0 | |

| | Full_meals_eaten | vitD_supp | Initial_days | Additional_charges |
|------|------------------|-----------|--------------|--------------------|
| mean | 0.0 | -0.0 | -0.0 | -0.0 |

```
In [166]: #Verifying the standard deviation of each column, should be 1 to confirm standardization
X_stand_stddv = X_stand_df.agg(['std']).round(2)

print(X_stand_stddv)
```

| | Lat | Lng | Population | Children | Age | Income | VitD_levels | Doc_visits | \ |
|-----|-----|-----|------------|----------|-----|--------|-------------|------------|---|
| std | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | |

| | Full_meals_eaten | vitD_supp | Initial_days | Additional_charges |
|-----|------------------|-----------|--------------|--------------------|
| std | 1.0 | 1.0 | 1.0 | 1.0 |

```
In [167]: # Provide a copy of the cleaned Data Set, index=False prevents the creation of an additional column  
X_stand_df.to_csv(r'C:\Users\mmorg\WGU\D212\Cleaned212Task2data.csv', index=False)
```

Part IV: Analysis

D. Perform PCA by doing the following:

D1. Determine the matrix of all the principal components.

Now that the data is standardized, we can use PCA to reduce dimensionality. We can use `fit_transform()` to combine both of these steps into one function.

After creating the matrix we can examine what weight each feature contributes to each of the principal components generated. For example PC1 is primarily based on `Initial_days`, whereas PC2 is influenced by `Longitude` and `Initial_days`. The higher the numbers on this matrix mean that those observations have a higher influence on those principle componets and we can start working on dimensionality reduction in the next step.

```
In [168]: # Instantiating data
pca = PCA(n_components = 12, random_state = 0)
# Fit and transform the PCA to X_stand
X_pca = pca.fit_transform(X_stand)
# Create PCA Matrix, shows the weight that a feature contributes to each Principal Component
X_pca_matrix = pd.DataFrame(pca.components_,
                           columns = ["PC1", "PC2", "PC3", "PC4", "PC5", "PC6",
                                      "PC7", "PC8", "PC9", "PC10", "PC11", "PC12"])
X_pca_matrix
```

Out[168]:

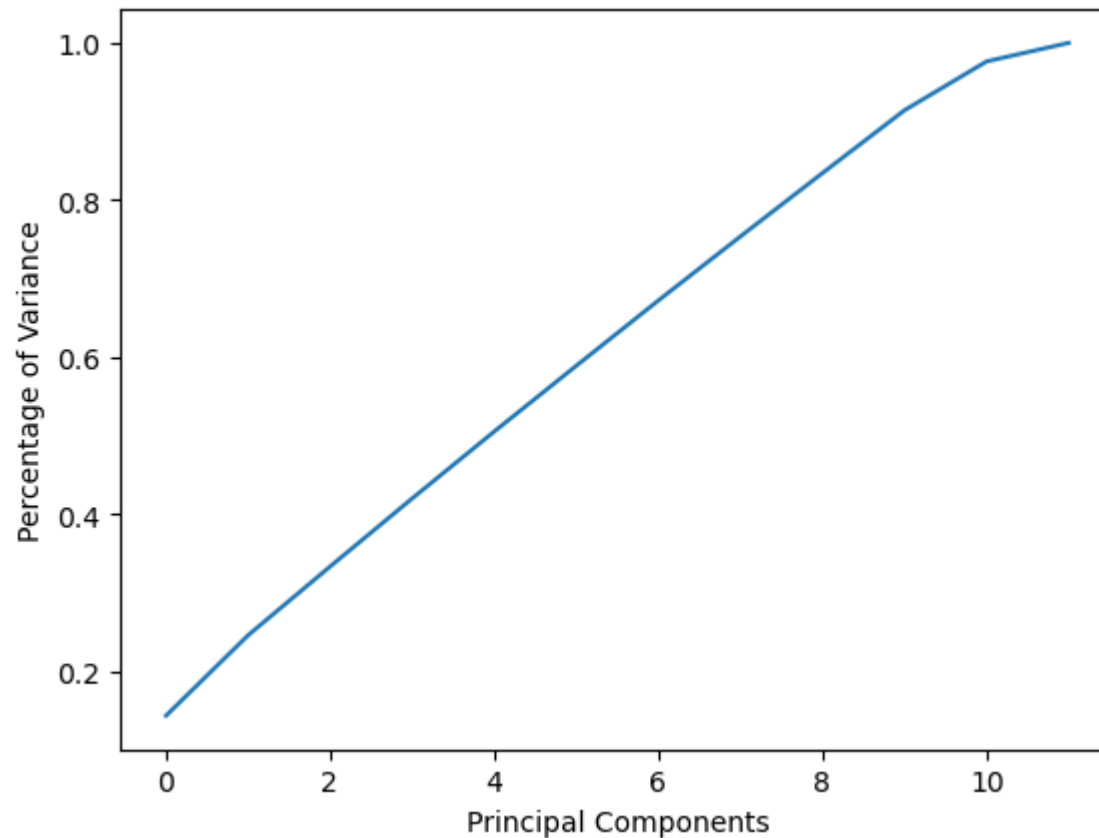
| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 | PC8 | PC9 | PC10 | PC11 | PC12 |
|----|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 | -0.003058 | 0.007652 | -0.022811 | 0.023338 | 0.705901 | -0.022410 | 0.019598 | 0.013893 | 0.027481 | 0.019038 | 0.020315 | 0.705664 |
| 1 | -0.712542 | 0.265927 | 0.629684 | -0.023013 | 0.008656 | 0.069845 | -0.021974 | 0.013776 | -0.111319 | 0.037972 | 0.070488 | 0.010325 |
| 2 | 0.129259 | -0.564164 | 0.253543 | 0.350595 | -0.003164 | 0.088519 | -0.080655 | 0.053630 | -0.299785 | 0.338816 | 0.505879 | -0.001590 |
| 3 | -0.064795 | -0.343004 | 0.234772 | 0.291952 | -0.022226 | -0.137444 | 0.560328 | 0.048673 | 0.531893 | -0.343571 | -0.032595 | -0.007786 |
| 4 | 0.041323 | -0.235515 | 0.102756 | -0.120302 | 0.009154 | 0.617901 | -0.011810 | 0.586993 | -0.028745 | 0.007754 | -0.435667 | 0.022723 |
| 5 | -0.042849 | 0.001443 | -0.052169 | 0.448856 | 0.013231 | 0.477300 | -0.391113 | -0.420891 | 0.033011 | -0.483289 | -0.030242 | 0.016962 |
| 6 | 0.007356 | -0.528847 | 0.273410 | -0.545779 | 0.017394 | -0.128018 | -0.297966 | -0.412969 | 0.118341 | -0.039959 | -0.234668 | 0.030821 |
| 7 | -0.057235 | 0.073339 | -0.067841 | 0.135775 | -0.014133 | 0.326654 | 0.193909 | -0.411312 | 0.351625 | 0.698886 | -0.206142 | -0.007117 |
| 8 | -0.052112 | 0.052695 | -0.008722 | 0.068395 | -0.020392 | -0.129952 | -0.592420 | 0.354429 | 0.655825 | 0.142106 | 0.213063 | -0.013102 |
| 9 | -0.002282 | 0.049267 | -0.111711 | -0.503772 | 0.006684 | 0.463827 | 0.219518 | -0.064247 | 0.196984 | -0.145438 | 0.635462 | -0.006324 |
| 10 | 0.679482 | 0.384103 | 0.614897 | -0.006411 | -0.001120 | 0.056180 | -0.003230 | -0.056510 | 0.074130 | -0.018554 | 0.003061 | 0.020499 |
| 11 | 0.008854 | -0.005042 | 0.016922 | 0.003751 | 0.707038 | 0.002297 | -0.002533 | 0.000685 | 0.010346 | 0.000585 | -0.011905 | -0.706705 |

D2. Identify the total number of principal components using the elbow rule or the Kaiser criterion.

Now that we have our 12 principal components we can start working on reducing those that have little influence. To do this we can create a plot and use the elbow rule. Below you can see the elbow plot and it starts to level out around 10 on the x-axis. Because it starts at 0, and the 0 is referring to PC1, this means that we want to keep 11 principal components heading forward.

The plot from this dataset isn't one of the most clear or definitive plots I've seen in previous coursework. However, there is a clear leveling out of the line after PC11, so we will drop our number of components down to 11 from here on out.

```
In [169]: # Use a scree plot to visualize the contribution of each PC to the whole of variance
# Scree plot
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel("Principal Components")
plt.ylabel("Percentage of Variance")
plt.show();
```



D3. Identify the variance of each of the principal components identified in part D2.

To identify the variance of each PCA identified in D2 we have to run the PCA again with only 11 components and produce a new matrix. Once we have the PCA created we can then assign the `explained_variance_ratio_` attribute. The numbers provided are percentage of variation of each principal component in decimal format.

```

In [170]: # Instantiating data
reduced_pca = PCA(n_components = 11, random_state = 0)
# Fit and transform the PCA to X_stand
reduced_pca.fit(X_stand)
final_pca = reduced_pca.transform(X_stand)
# Create PCA Matrix, shows the weight that a feature contributes to each Principal Component
final_pca_matrix = pd.DataFrame(reduced_pca.components_.T,
                                columns = ["PC1", "PC2", "PC3", "PC4", "PC5", "PC6",
                                           "PC7", "PC8", "PC9", "PC10", "PC11"])

final_pca_matrix

```

Out[170]:

| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 | PC8 | PC9 | PC10 | PC11 |
|----|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 | -0.003058 | -0.712542 | 0.129259 | -0.064795 | 0.041323 | -0.042849 | 0.007356 | -0.057235 | -0.052112 | -0.002282 | 0.679482 |
| 1 | 0.007652 | 0.265927 | -0.564164 | -0.343004 | -0.235515 | 0.001443 | -0.528847 | 0.073339 | 0.052695 | 0.049267 | 0.384103 |
| 2 | -0.022811 | 0.629684 | 0.253543 | 0.234772 | 0.102756 | -0.052169 | 0.273410 | -0.067841 | -0.008722 | -0.111711 | 0.614897 |
| 3 | 0.023338 | -0.023013 | 0.350595 | 0.291952 | -0.120302 | 0.448856 | -0.545779 | 0.135775 | 0.068395 | -0.503772 | -0.006411 |
| 4 | 0.705901 | 0.008656 | -0.003164 | -0.022226 | 0.009154 | 0.013231 | 0.017394 | -0.014133 | -0.020392 | 0.006684 | -0.001120 |
| 5 | -0.022410 | 0.069845 | 0.088519 | -0.137444 | 0.617901 | 0.477300 | -0.128018 | 0.326654 | -0.129952 | 0.463827 | 0.056180 |
| 6 | 0.019598 | -0.021974 | -0.080655 | 0.560328 | -0.011810 | -0.391113 | -0.297966 | 0.193909 | -0.592420 | 0.219518 | -0.003230 |
| 7 | 0.013893 | 0.013776 | 0.053630 | 0.048673 | 0.586993 | -0.420891 | -0.412969 | -0.411312 | 0.354429 | -0.064247 | -0.056510 |
| 8 | 0.027481 | -0.111319 | -0.299785 | 0.531893 | -0.028745 | 0.033011 | 0.118341 | 0.351625 | 0.655825 | 0.196984 | 0.074130 |
| 9 | 0.019038 | 0.037972 | 0.338816 | -0.343571 | 0.007754 | -0.483289 | -0.039959 | 0.698886 | 0.142106 | -0.145438 | -0.018554 |
| 10 | 0.020315 | 0.070488 | 0.505879 | -0.032595 | -0.435667 | -0.030242 | -0.234668 | -0.206142 | 0.213063 | 0.635462 | 0.003061 |
| 11 | 0.705664 | 0.010325 | -0.001590 | -0.007786 | 0.022723 | 0.016962 | 0.030821 | -0.007117 | -0.013102 | -0.006324 | 0.020499 |

```
In [171]: list(reduced_pca.explained_variance_ratio_)
```

```
Out[171]: [0.14329551589301778,  
          0.10240073568553085,  
          0.08757786111482174,  
          0.08667044956990438,  
          0.08521515256145173,  
          0.08373292555920939,  
          0.0831305698317373,  
          0.08150602610684196,  
          0.0807104238018291,  
          0.07998956523163969,  
          0.06221023077235041]
```

D4. Identify the total variance captured by the principal components identified in part D2.

Provided below is a running total of the variance of each principal component. This shows that the 11 principal components are responsible for 97.6% of total variance.

```
In [172]: list(reduced_pca.explained_variance_ratio_.cumsum())
```

```
Out[172]: [0.14329551589301778,  
          0.24569625157854863,  
          0.3332741126933704,  
          0.41994456226327476,  
          0.5051597148247264,  
          0.5888926403839359,  
          0.6720232102156731,  
          0.7535292363225151,  
          0.8342396601243443,  
          0.9142292253559839,  
          0.9764394561283344]
```

D5. Summarize the results of your data analysis.

Now that we've successfully performed our PCA, we can try to make some predictions and see how good of a job we did. We will now use a Decision Tree Classifier to try and classify patients based on our 11 remaining principal components and predict whether they will be re-admitted.

```
In [173]: # Create train and test sets from our final pca, which was already been fit and transformed
X_train, X_test, y_train, y_test = train_test_split(final_pca, y, train_size = 0.7, test_size=0.3, random_state=0)
# Confirm that our train and test datasets are shaped correctly
print(X_train.shape)
print(X_test.shape)
```

```
(7000, 11)
```

```
(3000, 11)
```

```
In [174]: # Instantiate DecisionTreeClassifier
classification_model = DecisionTreeClassifier(random_state=0).fit(X_train, y_train)
y_predicted = classification_model.predict(X_test)
# Check accuracy score
test_accuracy = accuracy_score(y_test, y_predicted)
print("The accuracy of the model")
print(test_accuracy)
# Predict probabilities of the positive class
y_predicted_probability = classification_model.predict_proba(X_test)[:,-1]
# Create confusion matrix
final_matrix = confusion_matrix(y_test, y_predicted)
# Print confusion matrix
print("\nThe confusion matrix for the Decision Tree Classifier")
print(final_matrix)
#Generate AUC score
print("\nThe Area Under the Curve (AUC) for the Decision Tree Classifier:")
#Print AUC score
print(roc_auc_score(y_test, y_predicted_probability))
```

```
The accuracy of the model
```

```
0.8873333333333333
```

```
The confusion matrix for the Decision Tree Classifier
```

```
[[1731  180]
```

```
 [ 158  931]]
```

```
The Area Under the Curve (AUC) for the Decision Tree Classifier:
```

```
0.880360620620361
```

After confirming that our train and test datasets are correct in shape we can continue forward with our model. Once we create the model we check the accuracy score and find that it is 88.7% accurate.

To see how it performs with raw numbers we create a confusion matrix. The read this matrix we have to understand what the four numbers mean. The left side numbers (1,731 and 158) refer to predictions of no re-admission. The right side numbers (180 and 931) refer to predictions of re-admission. The top row numbers (1,731 and 180) are actual patient results of no re-admission while the bottom row numbers (158 and 931) are actual patient results of re-admission.

So from reading this matrix we can see that of the 3,000 patients predicted, 2,662 (1,731 + 931) were predicted correctly while the remaining 338 (180 + 158) were incorrectly predicted.

Finally, our Area Under the Curve score can be looked at for even more evidence of our model's strength in predicting patient re-admissions. Our score here is 88% which is a great score and provides evidence that our model is strong.

In conclusion, I am somewhat shocked at these results. I have ran other prediction models on this dataset with similar accuracies. However, those analyses were using categorical data like health conditions and why the patient was hospitalized, in contrast to the continuous data of this dataset that for the most part seemingly should have no bearing on patient re-admission. Who would think that latitude and longitude could be used to predict patient re-admissions in this context? If there was something about where that person was living I would have to assume that it would result in one of the health conditions I mentioned as part of the categorical data analysis.

This definitely makes me think about a possible next step for this dataset and predicting patient re-admissions with even higher accuracy. If we could analyze this dataset with continuous and categorical data we could probably design a very accurate model that would outperform this model with an 89% accuracy score and 88% AUC score.

Part V: Attachments

E. Record the web sources used to acquire data or segments of third-party code to support the analysis. Ensure the web sources are reliable.

I used information from the Datacamp modules, and my previous submissions in D206 and D212 Task 1 to help complete the coding in this assignment.

F. Acknowledge sources, using in-text citations and references, for content that is quoted, paraphrased, or summarized.

Keita, Z. (2023). Principal Component Analysis in R Tutorial. <https://datacamp.com/tutorial/pca-analysis-r>

